

Buffer Pool Class

- getPage: fixed the method for the retrieve of a page (from disk) not present in the BufferPool. From first assignment the method was built to manage only HeapPages, now it must manage any kind of page including heap and betree pages (bug of my first assignment code).
- flushPage: from a given pid it takes the corresponding page in the BufferPool check if it's dirty and update the disk if dirty. Finally after updating the page in the disk it set the dirty bit to false.
- evictPage: the policy I built takes the first page in the bufferpool that has not been used. I extract the first page that is not dirty and then discard it. If in some rare case I have no pages in that state in the buffer Pool, I flush all the pages and delete the first of the list after the flush.
- discardPage: remove the page with the given pid.
- deleteTuple/insertTuple: I add or delete a tuple from a table. Both the method insertTuple() and deleteTuple() of the DBFile class returns the list of modified pages after the action of deleting or inserting. I need to update the buffer Pool for each modified page and mark the page as dirty. If the buffer pool is full I first call the evictPage method to make place for the new page.

BTreeFile Class

- findLeafPage: steps of the process
 - If the given page is a leaf page: I found the page I wanted; I can return it.
 - If the given page is an internal page: I still have to check the rules in order to go down to find the correct leaf page. I first check the special rule for field given == null and I go down to the left child. In all the other cases I pass through all the entries of the internal page and check: if the key is less than the page I want I go down to the left child; if there are no more entries next I go down to the right child. Every time I go down to the right or left child I recursively call the findleafpage method to check if I am still in another internal page or if I have found the leaf page.
- splitLeafPage/splitInternalPage: A first step is to move half of the elements (tuples for leafpage BTree entries for internal page) to a new page of same kind. Once done that process I need to update (only for leaf pages) the siblings pointers (in my logic I'm moving the elements from the first left part of the page so I have to change the pointers of the left sibling). For the internal page instead I have to delete the middle entry from the current page as in internal page we do not keep the element after moving it to the parent (another difference with splitting leaf page). Then I have to update the pointers and entries in the parent page. (comments in the code for more details).
 - ➔ Main difficulties: understand the correct pointer situation and how to update it.

Only 2 classes for the leaf pages for extra points

- stealFromLeafPage: I first check the number of pages to steal from the victim in order to maintain the correct balanced number of element in each page. Then with a left or right iterator, depending on the flag passed, I delete and add elements in order. Finally I update the parent element key that points to the 2 leaf page used, with the correct value that represent the new leaf page situation.
- mergeLeafPages: I delete all the elements from a victim page to move them to another page. I set the correct sibling (as I have deleted a page), mark the page as empty and delete the parent entry that was used to point to the 2 child pages now merged.

In the next page the high level schema of the simple db project 1 and 2. Drawing this schema helped me a lot to understand the connections between classes and build the methods for both projects.

Sources of informations: Javadoc, stackoverflow, github

